



Image Captioning Generator

Analysis of Methods for Optimizing a
Python-based Image Captioning Model

DS-GA 3001 Advanced Python

Charles Brillo-Sonnino
Meenakshi Jhalani

Aidan Claffey
Kevin Wilson

Darren Geng

May 5, 2020



- Dataset of 31,783 images and 158,915 sentence-based captions (5 per image)
- Task is to predict captions for previously unseen images, with many potential applications (e.g. accessibility, real-time video description)
- Many opportunities for optimization in data processing, model training, prediction and evaluation

Example image and associated captions

"An elderly man with light brown hair, wearing a gray sweater, blue shirt, brown pants and tan shoes, reads a book sitting at a bench, as two ducks feed themselves on the grass."

"A man with parted hair and wearing glasses is seated outdoors on a bench where he is reading."


"A man sits outside at a wooden table and reads a book while ducks eat in the foreground."

"An elderly man sitting on a bench while reading a book."

"Man reads in a park while feeding the ducks."



- Computing mean and standard deviation of array representation of RGB color channels across all images in the dataset (to normalize the data)
- Before optimizing, this operation took 472.9 seconds to run
- Chunking the data and using multithreading with N=4 processes shortened the time to 104.2 seconds



RGB Means: Multiprocessing with N=4 processes

```

filepath = 'flickr30k-images'
img_files = [filename for filename in os.listdir(filepath)]
chunks = [(img_files[i:i+500]) for i in range(0, len(img_files), 500)]

def getRGB(chunk):
    r_channel_sum = 0
    g_channel_sum = 0
    b_channel_sum = 0
    count = 0
    for filename in chunk:
        if filename[-3:] == 'jpg':
            img = np.array(Image.open(os.path.join(filepath, filename)).convert('RGB'))
            r_channel_sum += np.sum(img[:, :, 0])
            g_channel_sum += np.sum(img[:, :, 1])
            b_channel_sum += np.sum(img[:, :, 2])
            count += img.shape[0] * img.shape[1]
    return (r_channel_sum, g_channel_sum, b_channel_sum, count)

from multiprocessing.pool import Pool

start_time = time.time()
with Pool(4) as p:
    res = p.map(getRGB, chunks)

results = np.array(res).sum(axis=0)
r,g,b,c = results[0], results[1], results[2], results[3]

end_time = time.time()


print('R channel mean: {}'.format(r/c))
print('G channel mean: {}'.format(g/c))
print('B channel mean: {}'.format(b/c))
print("Time for Multiprocessing with N=4 processes: %secs" % (end_time - start_time))

R channel mean: 113.2971859326401
G channel mean: 107.42922106881713
B channel mean: 98.14465223794616
Time for Multiprocessing with N=4 processes: 47.823638916015625secs
    
```

Using Python Concurrency to Optimize Generation of Vocabulary

- Function `build_vocab` contains for loops which can be optimized for better performance
- Implementation of Multiprocessing with $N=4$ processes decreased the time from 26.2 seconds to 12.16 seconds

Generating Vocabulary Optimization with Concurrency



Multiprocessing with $N=4$ processes

```
def build_vocab(ann_file = '../flickr30k/results_20130124.token', threshold = 4):
    """Build a simple vocabulary wrapper."""
    punc_set = set(['.', ',', ';', ':', '!', '?', '(', ')'])
    counter = Counter()
    caption_list = []
    split = pickle.load(open('train_set.p', 'rb'))
    ann_file = os.path.expanduser(ann_file)
    with open(ann_file) as fh:
        for line in fh:
            img, caption = line.strip().split('\t')
            if img[-2] in split:
                caption_list.append(caption)

    pool = mp.Pool(4)
    tokens = pool.map(nltk.tokenize.word_tokenize, [caption.lower() for caption in tqdm(caption_list)])
    pool.close()
    tokens = [item for elem in tokens for item in elem]
    tokens = [elem for elem in tokens if elem not in punc_set]
    counter = Counter(tokens)

    # If the word frequency is less than 'threshold', then the word is discarded.
    words = [word for word, cnt in counter.items() if cnt >= threshold]

    # Create a vocab wrapper and add some special tokens.
    vocab = Vocabulary()
    vocab.add_word('<pad>')
    vocab.add_word('<start>')
    vocab.add_word('<end>')
    vocab.add_word('<unk>')
    vocab.add_word('<break>')

    # Add the words to the vocabulary.
    for i, word in enumerate(words):
        vocab.add_word(word)
    return vocab
```

```
start_time = time()

build_vocab()


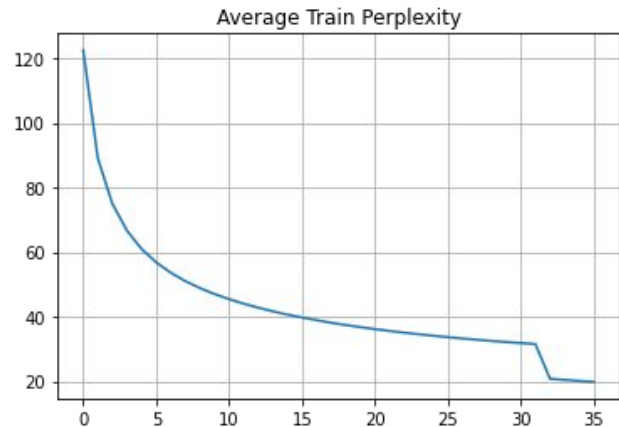
end_time = time()

print("Time for Multiprocessing with N=4 processes: %secs" % (end_time - start_time))

100% [#####] 141960/141960 [00:00<00:00, 875697.33it/s]

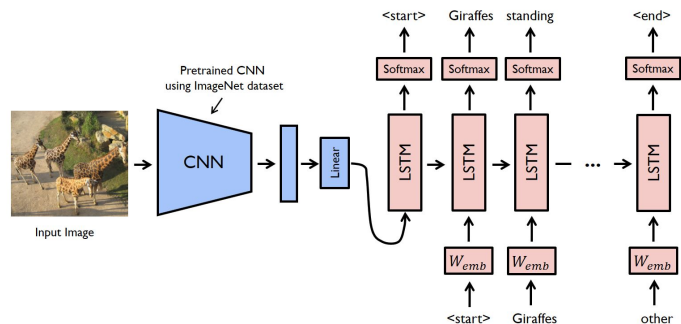
Time for Multiprocessing with N=4 processes: 12.161671161651611secs
```


Model Results : Training

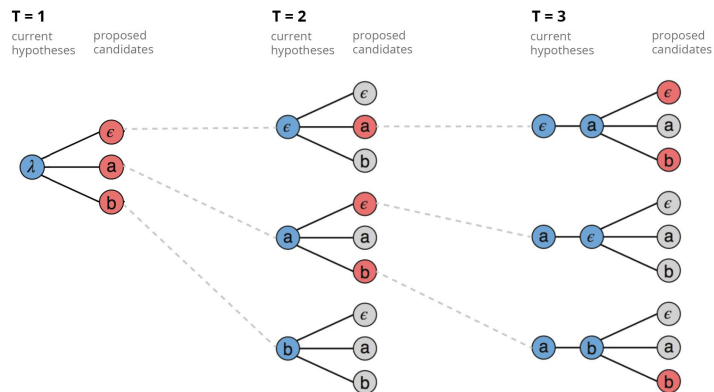


- *Encoder* : ResNet - 152 Pre-Trained on ImageNet
- *Decoder* : Single-layer LSTM with 512 hidden dimension and attention over image
- Training:
 - 30 epochs training just *Decoder*
 - 6 epoch training *Encoder* and *Decoder*
- Validation: Maximum BLEU-3 of **21.3**

Generating Text



- Decoder outputs the top-n best results at each timestep, along with their associated scores
- Greedy approach:
 - Always select the token with the best score
- Beam search:
 - Keep the best k results at each timestep, discard the rest



Standard beam search algorithm with an output alphabet $\{\epsilon, a, b\}$ and a beam size of three.

Timer unit: 1e-06 s

Total time: 1.30777 s

Function: beam_sample at line 200

Line #	Hits	Time	Per Hit	% Time	Line Contents
200					def beam_sample(self, features, targets=None, imgs=None, beam_size=3, max_seq_length=20, return_attention = False):
201					*** REMOVING IRRELEVANT CODE ***
201					"""Beam Search"""
236					
237	232	768.0	3.3	0.1	for k in range(beam_size):
238	2958	6736.0	2.3	0.5	for j in range(batch_size):
239	5568	12557.0	2.3	1.0	next_candidates[j].append(
240	5568	33011.0	5.9	2.5	(beam.scores[j].item() + topv[j][k].item(),
241	2784	18173.0	6.5	1.4	topi[j][k].item(),
242	2784	20555.0	7.4	1.6	beam.seq[j] + [str(topi[j][k].item())])
243)
244					
245					
246	2784	6841.0	2.5	0.5	if len(next_candidates[j]) > beam_size:
247	1824	5831.0	3.2	0.4	next_candidates[j].remove(min(next_candidates[j])) # only the top `beam_size` candidates are needed

Inefficient implementation

- Our first implementation continued several nested for loops and inefficient function call overhead
- Using line_profiler, we found that this part of the algorithm took about 0.1 seconds for 16 images (The main bottleneck of beam_search is in the encoding, which we could not optimize further)

Timer unit: 1e-06 s

Total time: 1.26072 s

Function: beam_sample at line 200

Line #	Hits	Time	Per Hit	% Time	Line Contents
200					def beam_sample(self, features, targets=None, imgs=None, beam_size=3, max_seq_length=20, return_attention = False):
201					"""Beam Search"""
*** REMOVING IRRELEVANT CODE ***					
256	986	2960.0	3.0	0.2	for j in range(batch_size):
257	1856	42397.0	22.8	3.4	next_candidates[j] += [(beam.scores[j].item()+topv[j][k],
258	928	2725.0	2.9	0.2	topi[j][k], beam.seq[j] + [str(topi[j][k])]) for k in range(beam_size)]
259					
260	20	1253.0	62.6	0.1	next_candidates = [sorted(next_cand)[-beam_size:] for next_cand in next_candidates]
*** REMOVING IRRELEVANT CODE ***					

```

def beam_sample(self, features, targets=None, imgs=None, beam_size=3, max_seq_length=20, return_attention = False):
    """Beam Search"""

*** REMOVING IRRELEVANT CODE ***

        for j in range(batch_size):
            next_candidates[j] += [(beam.scores[j].item()+topv[j][k],
                topi[j][k], beam.seq[j] + [str(topi[j][k])]) for k in range(beam_size)]

        next_candidates = [sorted(next_cand)[-beam_size:] for next_cand in next_candidates]

*** REMOVING IRRELEVANT CODE ***

```

Efficient implementation

- Updating the previous code with list-comprehension and fewer function calls allowed the program to run faster, taking up only approximately 0.05 seconds per 16 images, or an speed increase of 100%
- For a full validation set of 3000 images, this means a decrease in computation time of around 9 seconds



Greedy Search

a man in a black shirt and white hat is singing into a microphone while another man plays the drums

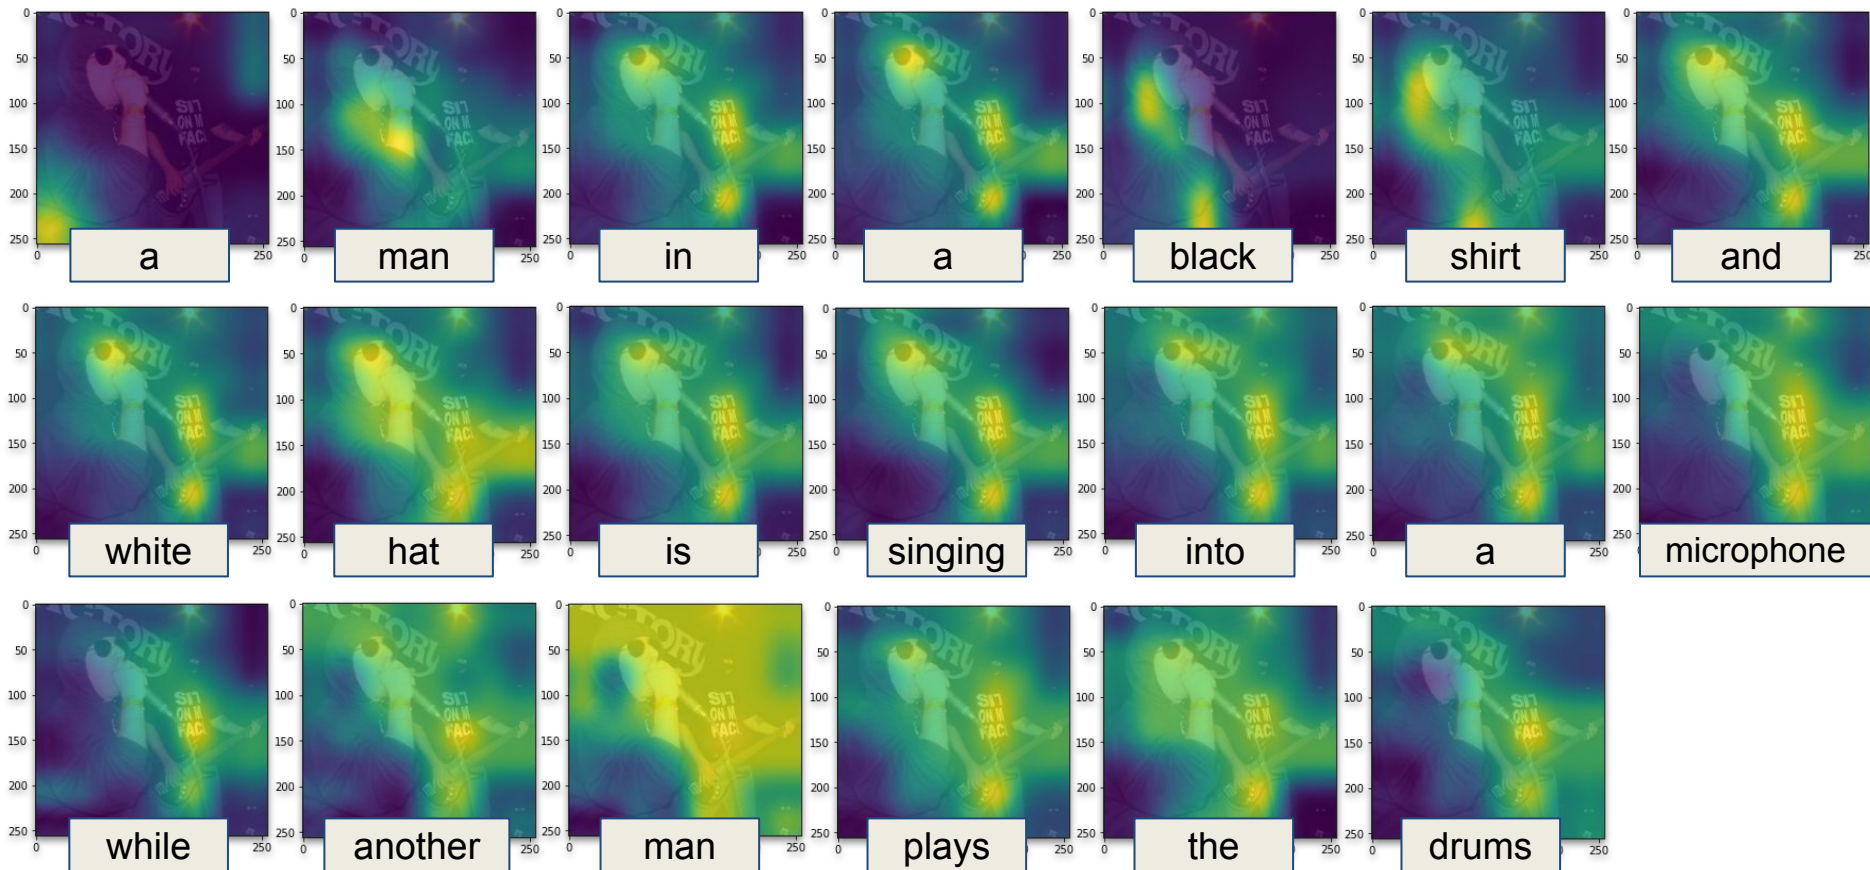
Beam = 3

man in white shirts and playing a flute with a man in a tattoo in the same with a beard

Beam = 5

three people are standing in the crowd behind them sandwich and the room with the sun shining a v belt

Model Results : Attention





Questions?